

7 How to anonymize columns

What will this tutorial cover?

In this tutorial we will cover six different techniques for anonymizing columns. We will find out how to use the function `fct_anon`, how to replace names with random names, how to mask values, how to group numeric variables, how to remove house numbers from street names, and how to encode and decode values.

Who do I have to thank?

For this course, I had to do some reading about anonymization. In particular, I relied on the following blog posts:

- <https://www.record-evolution.de/en/blog/data-anonymization-techniques-and-best-practices-a-quick-guide/>
- <https://satoricyber.com/data-masking/data-anonymization-use-cases-and-6-common-techniques/>
- <https://www.ssl2buy.com/wiki/symmetric-vs-asymmetric-encryption-what-are-differences>

No one wants to make their bank account password or home address public. When such data breaches occur, it is costly and potentially dangerous for those affected. For example in 2020, [private data of Finnish people in psychological care was hacked](#) and published. It even went so far that the hackers forced the victims to pay money to prevent their personal data from being disclosed.

To avoid such data breaches, we as users of R need to think twice about how we handle sensitive data. Therefore, in this tutorial we will cover some techniques to anonymize or pseudoanonymize your data. We will cover these topics:

- Anonymization of a factor column with `fct_anon`
- How to replace names with random names
- Masking values in a column
- Creating groups of numeric variables (e.g. age)
- Removing house numbers from street names
- Encrypting and decrypting columns with `encryptr`

What is the difference between pseudonymization and anonymization?

Before we begin, we need to talk about pseudonymization and anonymization. The following techniques can be used for pseudoanonymization and anonymization. The difference between the two is that pseudonymization is reversible, while anonymization is not (at least with current technical capabilities). All of the following techniques allow pseudonymization. Not all of them are suitable for anonymization (especially encryption).

The EU defines pseudonymization as follows:

“the processing of personal data in such a manner that the personal data can no longer be attributed to a specific data subject without the use of additional information provided that such additional information is kept separately and is subject to technical and organisational measures to ensure that the personal data are not attributed to an identified or identifiable natural person.” (https://edps.europa.eu/system/files/2021-04/21-04-27_aepd-edps_anonymisation_en_5.pdf)

By this definition, pseudonymization is reversible and requires additional information to reverse the process. Thus, for the following techniques, whether your data is anonymized or pseudoanonymized depends on your actions.

7.1 How to anonymize a factor column with `fct_anon`

Sometimes you want to make your data completely anonymous so that other people can't see sensitive information. For example, take a person's religion or gender. A simple function to anonymize such discrete data is `fct_anon`. The function takes two arguments. The factor you want to anonymize, and the prefix you put in front of the anonymized factor. Suppose we want to anonymize the factor levels of the `relig` column in the `gss_cat` data frame:

```
levels(gss_cat$relig)
```

```
[1] "No answer"          "Don't know"
[3] "Inter-nondenominational" "Native american"
[5] "Christian"          "Orthodox-christian"
[7] "Moslem/islam"       "Other eastern"
[9] "Hinduism"           "Buddhism"
[11] "Other"              "None"
[13] "Jewish"             "Catholic"
[15] "Protestant"         "Not applicable"
```

With `fct_anon` we can convert these levels into numeric values and add a prefix to them:

```
gss_cat %>%
  mutate(
    relig = fct_anon(relig, prefix = "religion_")
  ) %>%
  glimpse()
```

Rows: 21,483

Columns: 9

```
$ year      <int> 2000, 2000, 2000, 2000, 2000, 2000, 2000, 2000, 2000, 2000, 20~
$ marital   <fct> Never married, Divorced, Widowed, Never married, Divorced, Mar~
$ age       <int> 26, 48, 67, 39, 25, 25, 36, 44, 44, 47, 53, 52, 52, 51, 52, 40~
$ race      <fct> White, White, White, White, White, White, White, White, White, ~
$ rincome   <fct> $8000 to 9999, $8000 to 9999, Not applicable, Not applicable, ~
$ partyid   <fct> "Ind,near rep", "Not str republican", "Independent", "Ind,near~
$ relig     <fct> religion_02, religion_02, religion_02, religion_15, religion_1~
$ denom     <fct> "Southern baptist", "Baptist-dk which", "No denomination", "No~
$ tvhours   <int> 12, NA, 2, 4, 1, NA, 3, NA, 0, 3, 2, NA, 1, NA, 1, 7, NA, 3, 3~
```

Note that the numbers are generated randomly. So, each time you run this code, you will get a different set of numbers. If you need to use a factor as an id column, it is better to use other techniques that preserve the levels as unique identifiers (for example, hashing).

How to replace names with random names

Names are also sensitive data. To anonymize names, you can simply replace them with random names. This can be done with the `randomNames` function from the `randomNames` package:

```
library(randomNames)

presidential %>%
  mutate(
    name = randomNames(nrow(.),
                       sample.with.replacement = FALSE)
  )
```

A tibble: 11 x 4

| name | start | end | party |
|-------|--------|--------|-------|
| <chr> | <date> | <date> | <chr> |

| | | | | |
|----|---------------------|------------|------------|------------|
| 1 | al-Shaheen, Muammar | 1953-01-20 | 1961-01-20 | Republican |
| 2 | Stevenson, Daniel | 1961-01-20 | 1963-11-22 | Democratic |
| 3 | Salas, Zachary | 1963-11-22 | 1969-01-20 | Democratic |
| 4 | Burton, Charles | 1969-01-20 | 1974-08-09 | Republican |
| 5 | Nguyen, Joanne | 1974-08-09 | 1977-01-20 | Republican |
| 6 | Martinez, Gabrielle | 1977-01-20 | 1981-01-20 | Democratic |
| 7 | Whitetree, Sarina | 1981-01-20 | 1989-01-20 | Republican |
| 8 | Whittington, Ashtyn | 1989-01-20 | 1993-01-20 | Republican |
| 9 | Thompson, Brandon | 1993-01-20 | 2001-01-20 | Democratic |
| 10 | Velasquez, Veronica | 2001-01-20 | 2009-01-20 | Republican |
| 11 | Ahsan, Michael | 2009-01-20 | 2017-01-20 | Democratic |

In this example we have overwritten the column name with a list of random names and told the function that no name should occur more than once (`sample.with.replacement = FALSE`). Again, you get a different set of names each time you run the function.

If we want to be more specific about how the names are generated, we can provide some additional information to the function. For example, we can specify the ethnicity of the generated names, the order of their first and last names, and a separator that separates the first name from the last:

```
presidential %>%
  mutate(
    name = randomNames(nrow(.),
                       sample.with.replacement = FALSE,
                       ethnicity = c(1, 2, 4),
                       name.order = "first.last",
                       name.sep = " ")
  )
```

A tibble: 11 x 4

| | name <chr> | start <date> | end <date> | party <chr> |
|---|-----------------|-----------------|---------------|----------------|
| 1 | Atsila Yates | 1953-01-20 | 1961-01-20 | Republican |
| 2 | Yu Jin Sung | 1961-01-20 | 1963-11-22 | Democratic |
| 3 | Blanca Manchego | 1963-11-22 | 1969-01-20 | Democratic |
| 4 | William Duran | 1969-01-20 | 1974-08-09 | Republican |
| 5 | Grace Lam | 1974-08-09 | 1977-01-20 | Republican |
| 6 | Desirae Vasquez | 1977-01-20 | 1981-01-20 | Democratic |
| 7 | Ashley Palmier | 1981-01-20 | 1989-01-20 | Republican |
| 8 | Byunghyun Wang | 1989-01-20 | 1993-01-20 | Republican |
| 9 | Gerardo Meza | 1993-01-20 | 2001-01-20 | Democratic |

```
10 Matthew Zamora 2001-01-20 2009-01-20 Republican
11 Toshali Ngo    2009-01-20 2017-01-20 Democratic
```

If you are wondering what the names in the ethnicity vector mean, take a look at the documentation. Just hit `?randomNames` in your console and scroll down to ethnicity. For example, 1 stands for “American Indian or Native Alaskan”.

How to mask values (e.g. credit card numbers)

Another common use case is the masking of values. Credit card numbers are a prime example. You certainly don’t want to reveal a person’s credit card number. Masking is a technique that hides some characters of a string. Mostly by “X”s.

Suppose we want to mask the last digit of the height of Starwars characters. This is how we would do it in Tidyverse:

```
starwars %>%
  mutate(
    height = map_chr(height, ~ str_replace(.x, ".$", "X"))
  )
```

```
# A tibble: 87 x 14
   name      height mass hair_color skin_color eye_color birth_year sex  gender
   <chr>    <chr>  <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
1 Luke Sk~ 17X      77 blond     fair        blue        19   male masculi~
2 C-3PO    16X      75 <NA>      gold        yellow      112  none masculi~
3 R2-D2     9X      32 <NA>      white, bl~ red         33  none masculi~
4 Darth V~ 20X     136 none      white        yellow      41.9 male masculi~
5 Leia Or~ 15X      49 brown     light       brown        19  fema~ femin~
6 Owen La~ 17X     120 brown, gr~ light       blue         52  male masculi~
7 Beru Wh~ 16X      75 brown     light       blue         47  fema~ femin~
8 R5-D4     9X      32 <NA>      white, red red         NA   none masculi~
9 Biggs D~ 18X      84 black     light       brown        24  male masculi~
10 Obi-Wan~ 18X      77 auburn, w~ fair        blue-gray    57  male masculi~
# ... with 77 more rows, and 5 more variables: homeworld <chr>, species <chr>,
#   films <list>, vehicles <list>, starships <list>
```

The trick here lies in the `str_replace` function. The `.x` stands for the piped variable (in this case `height`). Then I provide a regular expression that searches for the last character of the string (`.$`). This character should then be replaced by an X.

A more complicated example arises when we want to mask more than one character. So let's return to our credit card example. Here is a data frame with credit card numbers:

```
ccards <- tibble(  
  creditcards = c(  
    36555224524299,  
    36350489667466,  
    36002887965170,  
    5447552069207504,  
    2221002654361034,  
    5127699386148536)  
)
```

Let's convert the first 10 characters of these credit card numbers to "X"s:

```
ccards %>%  
  mutate(  
    creditcards_masked = map_chr(creditcards, ~ str_replace(.x, "^.{10}",  
                                                             replacement = strrep("X", 10)))  
  )
```

```
# A tibble: 6 x 2  
  creditcards creditcards_masked  
    <dbl> <chr>  
1 3.66e13 XXXXXXXXXXXX4299  
2 3.64e13 XXXXXXXXXXXX7466  
3 3.60e13 XXXXXXXXXXXX5170  
4 5.45e15 XXXXXXXXXXXX207504  
5 2.22e15 XXXXXXXXXXXX361034  
6 5.13e15 XXXXXXXXXXXX148536
```

This code is a little more complicated than the first one. The regular expression `^.{10}` indicates that we are looking for the first 10 characters of the string. We replace this pattern with 10 "X"s, specified by `strrep("X", 10)`. The function `strrep` is a basic function of R, which simply repeats a series of characters:

```
strrep("X", 10)
```

```
[1] "XXXXXXXXXX"
```

Similarly, we could replace the last 5 characters with "X"s:

```
ccards %>%
  mutate(
    creditcars = map_chr(creditcards, ~ str_replace(.x, "\\d{5}$",
                                                    replacement = strrep("X", 5)))
  )
```

```
# A tibble: 6 x 2
  creditcards creditcars
      <dbl> <chr>
1 3.66e13 365552245XXXXX
2 3.64e13 363504896XXXXX
3 3.60e13 360028879XXXXX
4 5.45e15 54475520692XXXXX
5 2.22e15 22210026543XXXXX
6 5.13e15 51276993861XXXXX
```

In the regular expression `\\d{5}$` we look for the last five digits of a string.

How to turn ages into groups

Another common technique for anonymizing data is to divide it into groups. Suppose we want to divide the ages of people into groups.

Let's first calculate the age of our Starwars characters again:

```
(age_starwars <- starwars %>%
  mutate(age = as.integer(format(Sys.Date(), "%Y")) - birth_year) %>%
  select(name, age) %>%
  drop_na(age))
```

```
# A tibble: 43 x 2
  name          age
  <chr>        <dbl>
1 Luke Skywalker 2003
2 C-3PO         1910
3 R2-D2         1989
4 Darth Vader   1980.
5 Leia Organa   2003
6 Owen Lars     1970
7 Beru Whitesun lars 1975
```

```

8 Biggs Darklighter 1998
9 Obi-Wan Kenobi    1965
10 Anakin Skywalker 1980.
# ... with 33 more rows

```

Let us now go through four techniques we can use to group this age variable. With the function `cut_width` we can create groups of arbitrary width from a numeric variable. So let's group age into groups of 10 years:

```

age_starwars %>%
  mutate(
    age_groups = cut_width(age, 10)
  )

```

```

# A tibble: 43 x 3
  name          age age_groups
<chr>      <dbl> <fct>
1 Luke Skywalker 2003 (1995,2005]
2 C-3PO         1910 (1905,1915]
3 R2-D2         1989 (1985,1995]
4 Darth Vader   1980. (1975,1985]
5 Leia Organa   2003 (1995,2005]
6 Owen Lars     1970 (1965,1975]
7 Beru Whitesun lars 1975 (1965,1975]
8 Biggs Darklighter 1998 (1995,2005]
9 Obi-Wan Kenobi 1965 (1955,1965]
10 Anakin Skywalker 1980. (1975,1985]
# ... with 33 more rows

```

The round bracket means that a number is not included in the set. The square bracket means that a number is included in the set.

The function `cut_number` creates a certain number of sets. For example, we could say that the age column should be grouped into 10 groups:

```

age_starwars %>%
  mutate(
    age_groups = cut_number(age, 10)
  )

```



```
# A tibble: 43 x 3
  name          age age_groups
  <chr>         <dbl> <fct>
1 Luke Skywalker 2003 (2001,2014]
2 C-3PO         1910 [1126,1922]
3 R2-D2         1989 (1981,1991]
4 Darth Vader    1980. (1976,1981]
5 Leia Organa    2003 (2001,2014]
6 Owen Lars      1970 (1965,1970]
7 Beru Whitesun lars 1975 (1970,1976]
8 Biggs Darklighter 1998 (1991,2001]
9 Obi-Wan Kenobi  1965 (1965,1970]
10 Anakin Skywalker 1980. (1976,1981]
# ... with 33 more rows
```

Note, however, that the width of each group varies. For example, Luke Skywalker is in the set between 2011 and 2014 (a difference of 13 years) and Darth Vader is in the set between 1976 and 1981 (a difference of 5 years).

Then we can convert the age to millennia or decades. For example, Luke Skywalker is 2003 years old (at least in this data frame; I honestly don't know much about Starwars, so be gentle :)), so his age should fall in the millennia 2000. That's how we would do it with [David Robinson's](#) decade trick:

```
age_starwars %>%
  mutate(
    millenium = 1000 * (age %/% 1000)
  )
```

```
# A tibble: 43 x 3
  name          age millenium
  <chr>         <dbl>     <dbl>
1 Luke Skywalker 2003       2000
2 C-3PO         1910       1000
3 R2-D2         1989       1000
4 Darth Vader    1980.       1000
5 Leia Organa    2003       2000
6 Owen Lars      1970       1000
7 Beru Whitesun lars 1975       1000
8 Biggs Darklighter 1998       1000
9 Obi-Wan Kenobi  1965       1000
10 Anakin Skywalker 1980.       1000
```

```
# ... with 33 more rows
```

In a similar way, we could convert ages into decades:

```
age_starwars %>%  
  mutate(  
    decade = 10 * (age %/% 10)  
  )
```

```
# A tibble: 43 x 3
```

| | name | age | decade |
|----|--------------------|-------|--------|
| | <chr> | <dbl> | <dbl> |
| 1 | Luke Skywalker | 2003 | 2000 |
| 2 | C-3PO | 1910 | 1910 |
| 3 | R2-D2 | 1989 | 1980 |
| 4 | Darth Vader | 1980. | 1980 |
| 5 | Leia Organa | 2003 | 2000 |
| 6 | Owen Lars | 1970 | 1970 |
| 7 | Beru Whitesun lars | 1975 | 1970 |
| 8 | Biggs Darklighter | 1998 | 1990 |
| 9 | Obi-Wan Kenobi | 1965 | 1960 |
| 10 | Anakin Skywalker | 1980. | 1980 |

```
# ... with 33 more rows
```

How to remove house numbers from street names

Then we have street names and especially street numbers that sometimes need to be anonymized. Suppose we want to remove the street numbers of these streets:

```
street_names <- tibble(  
  street_name = c("Bromley Lanes 34",  
                  "Woodsgate Avenue 12",  
                  "Ardconnel Terrace 99",  
                  "Gipsy Birches 45",  
                  "Legate Close 8",  
                  "Stevenson Oval 9",  
                  "St Leonard's Boulevard 112",  
                  "Copper Chare 435",  
                  "Glastonbury Glebe 82",  
                  "Southern Way 91")  
)
```

To remove them, we can use the `str_remove_all` function from the `stringr` package:

```
street_names %>%  
  mutate(  
    street_names_no_number = str_remove_all(street_name, "\\d")  
  )
```

```
# A tibble: 10 x 2  
  street_name                street_names_no_number  
  <chr>                     <chr>  
1 Bromley Lanes 34          "Bromley Lanes "  
2 Woodsgate Avenue 12      "Woodsgate Avenue "  
3 Ardconnel Terrace 99     "Ardconnel Terrace "  
4 Gipsy Birches 45         "Gipsy Birches "  
5 Legate Close 8          "Legate Close "  
6 Stevenson Oval 9         "Stevenson Oval "  
7 St Leonard's Boulevard 112 "St Leonard's Boulevard "  
8 Copper Chare 435         "Copper Chare "  
9 Glastonbury Glebe 82     "Glastonbury Glebe "  
10 Southern Way 91         "Southern Way "
```

With the regular expression `\\d` we can remove all digits from a string.

How to encrypt and decrypt columns

Finally, we can anonymize each column by encrypting it. Encryption is a complicated and vast topic. I can't and won't go into detail here, but let me give you a brief introduction to how it works in R.

When we encrypt a column, we convert the values of a column into another form, which we call ciphertext. The ciphertext is not readable by humans, but it can be converted back to the original value. There are two forms of encryption. Symmetric encryption, where a single key is used to encrypt and decrypt a value, and asymmetric encryption, where two keys are used to encrypt and decrypt a value. A key is plaintext that translates between the two representations. Once you have the key in symmetric encryption, you can decrypt values. To decrypt values in asymmetric encryption, you need the public key and the private key.

The public key is as it says public, so open to anyone. The private key is a key you should not share with anyone. Only when you have both, can you decrypt a value. Also, private key cannot be guessed from the public key. To add another level of security, the private key also sometimes has a passphrase (or password) to it.

So much for the short theory, let's encrypt a column in R. For example, suppose you have this table of usernames and passwords:

```
users <- tibble(  
  name = c("Alexander", "Marie", "John"),  
  password = c(12345, "8$43_45*", "becker23#")  
)
```

We can encrypt this data with the package `encryptr`. First we need to load the package and create the private and public keys using the `genkeys` function:

```
library(encryptr)  
  
# genkeys() generates a public and private key pair  
# Passphrase: 456#7  
genkeys()
```

The function prompted us to provide a passphrase for the private key. This passphrase and the private key should not be shared with anyone!

Once we have the passphrase, we can encrypt our columns. Let's encrypt the password column:

```
users_encrypted <- users %>%  
  encrypt(password)  
  
users_encrypted %>%  
  glimpse()
```

```
Rows: 3  
Columns: 2  
$ name      <chr> "Alexander", "Marie", "John"  
$ password  <chr> "5b084bd19feef7b413b7a5ec4417f4c28b06e1f6ba8c91..."
```

As you can see, the passwords have been converted into a long string. It is not possible to decrypt the passwords from this string. To decrypt the column, we simply use the `decrypt` function:

```
users_encrypted %>% decrypt(password)
```

You must provide the passphrase to decrypt the column. Also, this works only if the R file is in the same directory as the public and private keys.

Summary

Here is what you can take from this tutorial.

- Anonymization and pseudoanonymization are two different techniques. Anonymization is not reversible, pseudoanonymization is.
- Whether your data is anonymized or pseudoanonymized depends on how you handle the data.
- With asymmetric encryption, you have a public key and a private key. Only with both keys you can decrypt a message.